

Kubectl Context and Configuration

Set which Kubernetes cluster kubectl communicates with and modifies configuration information. See [Authenticating Across Clusters with kubeconfig](#) documentation for detailed config file information.

```
kubectl config view # Show Merged kubeconfig settings.

# use multiple kubeconfig files at the same time and view merged config
KUBECONFIG=~/.kube/config:~/.kube/kubconfig2

kubectl config view

# get the password for the e2e user
kubectl config view -o jsonpath='{.users[?(@.name == "e2e")].user.password}'

kubectl config view -o jsonpath='{.users[].name}'      # display the first user
kubectl config view -o jsonpath='{.users[*].name}'    # get a list of users
kubectl config get-contexts                          # display list of contexts
kubectl config current-context                      # display the current-
context
kubectl config use-context my-cluster-name          # set the default context to
my-cluster-name

# add a new user to your kubeconf that supports basic auth
kubectl config set-credentials kubeuser/foo.kubernetes.com --username=kubeuser --
password=kubepassword

# permanently save the namespace for all subsequent kubectl commands in that
context.
kubectl config set-context --current --namespace=ggckad-s2

# set a context utilizing a specific username and namespace.
kubectl config set-context gce --user=cluster-admin --namespace=foo \
  && kubectl config use-context gce

kubectl config unset users.foo                    # delete user foo
```

Creating Objects

Kubernetes manifests can be defined in YAML or JSON. The file extension .yaml, .yml, and .json can be used.

```
kubectl apply -f ./my-manifest.yaml      # create resource(s)
kubectl apply -f ./my1.yaml -f ./my2.yaml # create from multiple files
kubectl apply -f ./dir                  # create resource(s) in all
manifest files in dir
```

```

kubect1 apply -f https://git.io/vPieo           # create resource(s) from url
kubect1 create deployment nginx --image=nginx   # start a single instance of nginx
kubect1 explain pods                           # get the documentation for pod
manifests

# Create multiple YAML objects from stdin
cat <<EOF | kubect1 apply -f -
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep
spec:
  containers:
  - name: busybox
    image: busybox
    args:
    - sleep
    - "1000000"
---
apiVersion: v1
kind: Pod
metadata:
  name: busybox-sleep-less
spec:
  containers:
  - name: busybox
    image: busybox
    args:
    - sleep
    - "1000"
EOF

# Create a secret with several keys
cat <<EOF | kubect1 apply -f -
apiVersion: v1
kind: Secret
metadata:
  name: mysecret
type: Opaque
data:
  password: $(echo -n "s33msi4" | base64 -w0)
  username: $(echo -n "jane" | base64 -w0)
EOF

```

Viewing, Finding Resources

```

# Get commands with basic output
kubect1 get services           # List all services in the
namespace                     # List all pods in all namespaces
kubect1 get pods --all-namespaces

```

```

kubect1 get pods -o wide                    # List all pods in the current
namespace, with more details
kubect1 get deployment my-dep              # List a particular deployment
kubect1 get pods                          # List all pods in the namespace
kubect1 get pod my-pod -o yaml             # Get a pod's YAML

# Describe commands with verbose output
kubect1 describe nodes my-node
kubect1 describe pods my-pod

# List Services Sorted by Name
kubect1 get services --sort-by=.metadata.name

# List pods Sorted by Restart Count
kubect1 get pods --sort-by='.status.containerStatuses[0].restartCount'

# List PersistentVolumes sorted by capacity
kubect1 get pv --sort-by=.spec.capacity.storage

# Get the version label of all pods with Label app=cassandra
kubect1 get pods --selector=app=cassandra -o \
  jsonpath='{.items[*].metadata.labels.version}'

# Retrieve the value of a key with dots, e.g. 'ca.crt'
kubect1 get configmap myconfig \
  -o jsonpath='{.data.ca\.crt}'

# Get all worker nodes (use a selector to exclude results that have a label
# named 'node-role.kubernetes.io/master')
kubect1 get node --selector='!node-role.kubernetes.io/master'

# Get all running pods in the namespace
kubect1 get pods --field-selector=status.phase=Running

# Get ExternalIPs of all nodes
kubect1 get nodes -o
jsonpath='{.items[*].status.addresses[?(@.type=="ExternalIP")].address}'

# List Names of Pods that belong to Particular RC
# "jq" command useful for transformations that are too complex for jsonpath, it
can be found at https://stedolan.github.io/jq/
sel=${$(kubect1 get rc my-rc --output=json | jq -j '.spec.selector | to_entries |
.[] | "\(.key)=\(.value),"'%?)}
echo ${$(kubect1 get pods --selector=$sel --
output=jsonpath='{.items..metadata.name}')}

# Show labels for all pods (or any other Kubernetes object that supports
labelling)
kubect1 get pods --show-labels

# Check which nodes are ready

```

```

JSONPATH='{range .items[*]}{@.metadata.name}:{range
@.status.conditions[*]}{@.type}={@.status};{end}{end}}' \
&& kubectl get nodes -o jsonpath="$JSONPATH" | grep "Ready=True"

# List all Secrets currently in use by a pod
kubectl get pods -o json | jq
'.items[].spec.containers[].env[]?.valueFrom.secretKeyRef.name' | grep -v null |
sort | uniq

# List all containerIDs of initContainer of all pods
# Helpful when cleaning up stopped containers, while avoiding removal of
initContainers.
kubectl get pods --all-namespaces -o jsonpath='{range
.items[*].status.initContainerStatuses[*]}{.containerID}{"\n"}{end}}' | cut -d/ -
f3

# List Events sorted by timestamp
kubectl get events --sort-by=.metadata.creationTimestamp

# Compares the current state of the cluster against the state that the cluster
would be in if the manifest was applied.
kubectl diff -f ./my-manifest.yaml

```

Updating Resources

```

kubectl set image deployment/frontend www=image:v2           # Rolling update
"www" containers of "frontend" deployment, updating the image
kubectl rollout history deployment/frontend                   # Check the
history of deployments including the revision
kubectl rollout undo deployment/frontend                       # Rollback to
the previous deployment
kubectl rollout undo deployment/frontend --to-revision=2      # Rollback to a
specific revision
kubectl rollout status -w deployment/frontend                 # Watch rolling
update status of "frontend" deployment until completion
kubectl rollout restart deployment/frontend                   # Rolling
restart of the "frontend" deployment

cat pod.json | kubectl replace -f -                           # Replace a pod
based on the JSON passed into std

# Force replace, delete and then re-create the resource. Will cause a service
outage.
kubectl replace --force -f ./pod.json

# Create a service for a replicated nginx, which serves on port 80 and connects
to the containers on port 8000
kubectl expose rc nginx --port=80 --target-port=8000

```

```

# Update a single-container pod's image version (tag) to v4
kubectl get pod mypod -o yaml | sed 's/\(image: myimage\):.*$/\1:v4/' | kubectl
replace -f -

kubectl label pods my-pod new-label=awesome # Add a Label
kubectl annotate pods my-pod icon-url=http://goo.gl/XXBTWq # Add an
annotation
kubectl autoscale deployment foo --min=2 --max=10 # Auto scale a
deployment "foo"

```

Patching Resources

```

# Partially update a node
kubectl patch node k8s-node-1 -p '{"spec":{"unschedulable":true}}'

# Update a container's image; spec.containers[*].name is required because it's a
merge key
kubectl patch pod valid-pod -p '{"spec":{"containers":[{"name":"kubernetes-serve-
hostname","image":"new image"}]}}'

# Update a container's image using a json patch with positional arrays
kubectl patch pod valid-pod --type='json' -p='[{"op": "replace", "path":
"/spec/containers/0/image", "value":"new image"}]'

# Disable a deployment livenessProbe using a json patch with positional arrays
kubectl patch deployment valid-deployment --type json -p='[{"op": "remove",
"path": "/spec/template/spec/containers/0/livenessProbe"}]'

# Add a new element to a positional array
kubectl patch sa default --type='json' -p='[{"op": "add", "path": "/secrets/1",
"value": {"name": "whatever" } }]'

```

Editing Resources

Edit any API resource in your preferred editor.

```

kubectl edit svc/docker-registry # Edit the service named
docker-registry
KUBE_EDITOR="nano" kubectl edit svc/docker-registry # Use an alternative editor

```

Scaling Resources

```

kubect1 scale --replicas=3 rs/foo                                # Scale a
replicaset named 'foo' to 3
kubect1 scale --replicas=3 -f foo.yaml                          # Scale a
resource specified in "foo.yaml" to 3
kubect1 scale --current-replicas=2 --replicas=3 deployment/mysql # If the
deployment named mysql's current size is 2, scale mysql to 3
kubect1 scale --replicas=5 rc/foo rc/bar rc/baz                 # Scale
multiple replication controllers

```

Deleting Resources

```

kubect1 delete -f ./pod.json                                    #
Delete a pod using the type and name specified in pod.json
kubect1 delete pod,service baz foo                             #
Delete pods and services with same names "baz" and "foo"
kubect1 delete pods,services -l name=myLabel                   #
Delete pods and services with label name=myLabel
kubect1 -n my-ns delete pod,svc --all                           #
Delete all pods and services in namespace my-ns,
# Delete all pods matching the awk pattern1 or pattern2
kubect1 get pods -n mynamespace --no-headers=true | awk
'/pattern1|pattern2/{print $1}' | xargs kubect1 delete -n mynamespace pod

```

Interacting with running Pods

```

kubect1 logs my-pod                                             # dump pod logs (stdout)
kubect1 logs -l name=myLabel                                    # dump pod logs, with label
name=myLabel (stdout)
kubect1 logs my-pod --previous                                  # dump pod logs (stdout) for
a previous instantiation of a container
kubect1 logs my-pod -c my-container                             # dump pod container logs
(stdout, multi-container case)
kubect1 logs -l name=myLabel -c my-container                   # dump pod logs, with label
name=myLabel (stdout)
kubect1 logs my-pod -c my-container --previous                 # dump pod container logs
(stdout, multi-container case) for a previous instantiation of a container
kubect1 logs -f my-pod                                          # stream pod logs (stdout)
kubect1 logs -f my-pod -c my-container                         # stream pod container logs
(stdout, multi-container case)
kubect1 logs -f -l name=myLabel --all-containers               # stream all pods logs with
label name=myLabel (stdout)
kubect1 run -i --tty busybox --image=busybox -- sh            # Run pod as interactive
shell
kubect1 run nginx --image=nginx -n                             # Run pod nginx in a specific
mynamespace namespace
kubect1 run nginx --image=nginx                                # Run pod nginx and write its
spec into a file called pod.yaml
--dry-run=client -o yaml > pod.yaml

```

```

kubect1 attach my-pod -i                # Attach to Running Container
kubect1 port-forward my-pod 5000:6000  # Listen on port 5000 on the
Local machine and forward to port 6000 on my-pod
kubect1 exec my-pod -- ls /             # Run command in existing pod
(1 container case)
kubect1 exec my-pod -c my-container -- ls / # Run command in existing pod
(multi-container case)
kubect1 top pod POD_NAME --containers  # Show metrics for a given
pod and its containers

```

Interacting with Nodes and Cluster

```

kubect1 cordon my-node                  # Mark my-
node as unschedulable
kubect1 drain my-node                  # Drain my-
node in preparation for maintenance
kubect1 uncordon my-node               # Mark my-
node as schedulable
kubect1 top node my-node               # Show
metrics for a given node
kubect1 cluster-info                   # Display
addresses of the master and services
kubect1 cluster-info dump              # Dump
current cluster state to stdout
kubect1 cluster-info dump --output-directory=/path/to/cluster-state # Dump
current cluster state to /path/to/cluster-state

# If a taint with that key and effect already exists, its value is replaced as
specified.
kubect1 taint nodes foo dedicated=special-user:NoSchedule

```

Resource types

List all supported resource types along with their shortnames, [API group](#), whether they are [namespaced](#), and [Kind](#):

```
kubect1 api-resources
```

Other operations for exploring API resources:

```

kubect1 api-resources --namespaced=true # ALL namespaced resources
kubect1 api-resources --namespaced=false # ALL non-namespaced resources
kubect1 api-resources -o name           # ALL resources with simple output
(just the resource name)
kubect1 api-resources -o wide           # ALL resources with expanded (aka
"wide") output
kubect1 api-resources --verbs=list,get  # ALL resources that support the
"list" and "get" request verbs
kubect1 api-resources --api-group=extensions # ALL resources in the "extensions"
API group

```

Formatting output

To output details to your terminal window in a specific format, add the `-o` (or `--output`) flag to a supported `kubectl` command.

Output format	Description
<code>-o=custom-columns=<spec></code>	Print a table using a comma separated list of custom columns
<code>-o=custom-columns-file=<filename></code>	Print a table using the custom columns template in the <code><filename></code> file
<code>-o=json</code>	Output a JSON formatted API object
<code>-o=jsonpath=<template></code>	Print the fields defined in a jsonpath expression
<code>-o=jsonpath-file=<filename></code>	Print the fields defined by the jsonpath expression in the <code><filename></code> file
<code>-o=name</code>	Print only the resource name and nothing else
<code>-o=wide</code>	Output in the plain-text format with any additional information, and for pods, the node name is included
<code>-o=yaml</code>	Output a YAML formatted API object

Examples using `-o=custom-columns`:

```
# All images running in a cluster
kubectl get pods -A -o=custom-columns='DATA:spec.containers[*].image'

# All images excluding "k8s.gcr.io/coredns:1.6.2"
kubectl get pods -A -o=custom-columns='DATA:spec.containers[?(!.image!="k8s.gcr.io/coredns:1.6.2")].image'

# All fields under metadata regardless of name
kubectl get pods -A -o=custom-columns='DATA:metadata.*'
```

Kubectl output verbosity and debugging

Kubectl verbosity is controlled with the `-v` or `--v` flags followed by an integer representing the log level. General Kubernetes logging conventions and the associated log levels are described [here](#).

Verbosity	Description
<code>--v=0</code>	Generally useful for this to <i>always</i> be visible to a cluster operator.
<code>--v=1</code>	A reasonable default log level if you don't want verbosity.

Verbosity	Description
--v=2	Useful steady state information about the service and important log messages that may correlate to significant changes in the system. This is the recommended default log level for most systems.
--v=3	Extended information about changes.
--v=4	Debug level verbosity.
--v=6	Display requested resources.
--v=7	Display HTTP request headers.
--v=8	Display HTTP request contents.
--v=9	Display HTTP request contents without truncation of contents.